

CS5740: Natural Language Processing
Spring 2016

IBM Translation Models

Instructor: Yoav Artzi

Slides adapted from Michael Collins

The Noisy Channel Model

- Goal: **translate from French to English**
- Have a model $p(e|f)$ to estimate the probability of an English sentence e given a French sentence f
- Estimate the parameters from training corpus
- A noisy channel model has two components:

$p(e)$ **the language model**
 $p(f|e)$ **the translation model**

- Giving:

$$p(e|f) = \frac{p(e, f)}{p(f)} = \frac{p(e)p(f|e)}{\sum_e p(e)p(f|e)}$$

and

$$\arg \max_e p(e|f) = \arg \max_e p(e)p(f|e)$$

Overview

- IBM Model 1
- IBM Model 2
- EM Training of Models 1 and 2

IBM Model 1: Alignments

- How do we model $p(f|e)$?
- English sentence e has l words $e^1 \dots e^l$
French sentence f has m words $f^1 \dots f^m$
- An **alignment** a identifies which English word each French word originated from
- Formally, an alignment a is:
$$\{a_1, \dots, a_m\} \quad \text{where } a_j \in 0 \dots l$$
- There are $(l + 1)^m$ possible alignments



IBM Model 1: Alignments

$l = 6, m = 7$

e = And the program has been implemented

f = Le programme a ete mis en application

IBM Model 1: Alignments

$l = 6, m = 7$

e = And the program has been implemented

f = Le programme a ete mis en application



- One alignment is

$\{2, 3, 4, 5, 6, 6, 6\}$

IBM Model 1: Alignments

$l = 6, m = 7$

e = And the program has been implemented

f = Le programme a ete mis en application

- Another (bad!) alignment is


$\{1, 1, 1, 1, 1, 1, 1\}$

IBM Model 1: Alignments

$l = 6, m = 7$

$e =$ And the program has been implemented

$f =$ Le programme a ete mis en application



- Another (bad!) alignment is

$\{1, 1, 1, 1, 1, 1, 1\}$

Alignments in the IBM Models

- We define two models:

$$\underline{p(a|e, m)} \quad \underline{p(f|a, e, m)}$$

- Giving:

$$p(f, a|e, m) = \underline{p(a|e, m)} \underline{p(f|a, e, m)}$$

- Also:

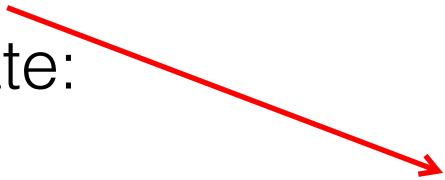
$$\underline{p(f|e, m)} = \sum_{a \in \mathcal{A}} \underline{p(a|e, m)} \underline{p(f|a, e, m)}$$

where \mathcal{A} is a set of all possible alignments

Most Likely Alignments

$$p(f, a|e, m) = p(a|e, m)p(f|a, e, m)$$

- We can also calculate:


$$p(a|f, e, m) = \frac{p(f, a|e, m)}{\sum_{a \in \mathcal{A}} p(f, a|e, m)}$$

for any alignment a

- For a given f, e pair, can also compute the most likely alignment (details in notes)
- The original IBM models are rarely used for translation, but still key for recovering alignments

Example Alignment

- French:
le conseil a rendu son avis , et nous devons à présent adopter un
nouvel avis sur la base de la première position .
- English:
the council has stated its position , and now , on the basis of the
first position , we again have to give our opinion .
- Alignment:
the/**le** council/**conseil** has/**à** stated/**rendu** its/**son** position/**avis** ,/
and/**et** now/**présent** ,/**NULL** on/**sur** the/**le** basis/**base** of/**de** the/**la**
first/**première** position/**position** ,/**NULL** we/**nous** again/**NULL**
have/**devons** to/**a** give/**adopter** our/**nouvel** opinion/**avis** ./.

IBM Model 1: Alignments

- In IBM Model 1 all alignments a are equally likely:

$$\underline{p(a|e, m)} = \frac{1}{(1 + l)^m}$$

- Reasonable assumption?
 - Simplifying assumption, but it gets things started ...

IBM Model 1: Translation Probabilities

- Next step: come up with an estimate for

$$\underline{p(f|a, e, m)}$$

- In Model 1, this is:

$$\underline{p(f|a, e, m)} = \prod_{j=1}^m t(f_j | e_{a_j})$$



IBM Model 1: Example

$l = 6, m = 7$

e = And the program has been implemented

f = Le programme a ete mis en application

$a = \{2, 3, 4, 5, 6, 6, 6\}$

IBM Model 1: Example

p(file)	And	the	program	has	been	implemented
Le	0.2	0.6	0.1	0.025	0.05	0.025
programme	0.05	0.2	0.45	0.1	0.1	0.1
a	0.1	0.1	0.15	0.2	0.15	0.3
ete	0.05	0.05	0.05	0.05	0.7	0.1
mis	0.2	0.05	0.05	0.05	0.25	0.4
en	0.25	0.1	0.25	0.25	0.1	0.05
application	0.01	0.03	0.01	0.02	0.03	0.9

IBM Model 1: Example

$$l = 6, m = 7$$

e = And the program has been implemented

f = Le programme a ete mis en application

$$a = \{2, 3, 4, 5, 6, 6, 6\}$$

$$\begin{aligned} p(f|a, e) = & t(\text{Le}|\text{the}) \times t(\text{programme}|\text{program}) \\ & \times t(\text{a}|\text{has}) \times t(\text{ete}|\text{been}) \\ & \times t(\text{mis}|\text{implemented}) \times t(\text{en}|\text{implemented}) \\ & \times t(\text{application}|\text{implemented}) = 0.0006804 \end{aligned}$$

$$p(f, a \mid e, 7) = 8.26186E - 10$$

IBM Model 1: The Generative Process

Translating
French → English

To generate a French string f from an English string e :

- Step 1: Pick an alignment a with probability $\frac{1}{(l+1)^m}$
- Step 2: Pick the French words with probability

$$\underline{p(f|a, e, m)} = \prod_{j=1}^m t(f_j | e_{a_j})$$

The final result:

$$\underline{p(f, a|e, m)} = \underline{p(a|e, m)} \times \underline{p(f|a, e, m)} = \frac{1}{(1+l)^m} \prod_{j=1}^m t(f_j | e_{a_j})$$

Example Lexical Entry

English	French	Probability
position	position	0.756715
position	situation	0.0547918
position	mesure	0.0281663
position	vue	0.0169303
position	point	0.0124795
position	attitude	0.0108907

... de la **situation** au niveau des négociations de l'ompi ...
... of the current **position** in the wipo negotiations ...

nous ne sommes pas en **mesure** de décider, ...
we are not in **position** to decide ...

... Le **point de vue** de la commission face à ce problème complexe .
... the commission 's **position** on this complex problem .

Overview

- IBM Model 1
- IBM Model 2
- EM Training of Models 1 and 2

IBM Model 2

- Only difference: we now introduce **alignment distortion** parameters

$$q(i|j, l, m)$$

- Probability that j 'th French word is connected to i 'th English word, given sentence length of e and f are l and m
- Define

$$p(a|e, m) = \prod_{j=1}^m q(a_j|j, l, m)$$

where $a = \{a_1, \dots, a_m\}$

- Gives

$$p(f, a|e, m) = \prod_{j=1}^m q(a_j|j, l, m) t(f_j|e_{a_j})$$



Example

$l = 6$

$m = 7$

$e =$ And the program has been implemented

$f =$ Le programme a ete mis en application

$a = \{2, 3, 4, 5, 6, 6, 6\}$

Example

$$l = 6$$

$$m = 7$$

$$e = \text{And the program has been implemented}$$

$$f = \text{Le programme a ete mis en application}$$

$$a = \{2, 3, 4, 5, 6, 6, 6\}$$

$$\begin{aligned} p(a \mid e, 7) = & \mathbf{q}(2 \mid 1, 6, 7) \times \\ & \mathbf{q}(3 \mid 2, 6, 7) \times \\ & \mathbf{q}(4 \mid 3, 6, 7) \times \\ & \mathbf{q}(5 \mid 4, 6, 7) \times \\ & \mathbf{q}(6 \mid 5, 6, 7) \times \\ & \mathbf{q}(6 \mid 6, 6, 7) \times \\ & \mathbf{q}(6 \mid 7, 6, 7) \end{aligned}$$

Example

$$l = 6$$

$$m = 7$$

$$e = \text{And the program has been implemented}$$

$$f = \text{Le programme a ete mis en application}$$

$$a = \{2, 3, 4, 5, 6, 6, 6\}$$

$$\begin{aligned} p(f \mid a, e, 7) = & \mathbf{t}(Le \mid the) \times \\ & \mathbf{t}(programme \mid program) \times \\ & \mathbf{t}(a \mid has) \times \\ & \mathbf{t}(ete \mid been) \times \\ & \mathbf{t}(mis \mid implemented) \times \\ & \mathbf{t}(en \mid implemented) \times \\ & \mathbf{t}(application \mid implemented) \end{aligned}$$

IBM Model 2: The Generative Process

Translating
French → English

To generate a French string f from an English string e :

- Step 1: Pick an alignment $a = \{a_1, \dots, a_m\}$ with probability
- Step 2: Pick the French words with probability

$$p(a|e, m) = \prod_{j=1}^m q(a_j | j, l, m)$$

$$p(f|a, e, m) = \prod_{j=1}^m t(f_j | e_{a_j})$$

The final result:

$$p(f, a|e, m) = p(a|e, m) \times p(f|a, e, m) = \prod_{j=1}^m q(a_j | j, l, m) t(f_j | e_{a_j})$$

Recovering Alignments

- If we have parameters q and t , we can easily recover the most likely alignment for any sentence pair

Given a sentence pair

$$e_1, e_2, \dots, e_l, f_1, f_2, \dots, f_m$$

define

$$a_j = \arg \max_{a \in \{0 \dots l\}} q(a|j, l, m) \times t(f_j, e_a)$$

for $j = 1 \dots m$

e = And the program has been implemented

f = Le programme a ete mis en application

Overview

- IBM Model 1
- IBM Model 2
- EM Training of Models 1 and 2

The Parameter Estimation Problem

- Input:

$$(e^{(k)}, f^{(k)}), k = 1 \dots n$$

Each $e^{(k)}$ is an English sentence, each $f^{(k)}$ is a French sentence

- Output: parameter for

$$t(f|e) \quad q(i|j, l, m)$$

- A key challenge: we do not have alignments in our training examples

$e^{(100)}$ = And the program has been implemented

$f^{(100)}$ = Le programme a ete mis en application

Parameter Estimation if Alignments are Observed

- Assume alignments are observed in training data
 $e^{(100)} =$ And the program has been implemented

$f^{(100)} =$ Le programme a ete mis en application

$a^{(100)} = \langle 2, 3, 4, 5, 6, 6, 6 \rangle$

- Training data is

$$(e^{(k)}, f^{(k)}, a^{(k)}), k = 1 \dots n$$

Each $e^{(k)}$ is an English sentence, each $f^{(k)}$ is a French sentence, each $a^{(k)}$ is an alignment

- Maximum-likelihood parameter estimates are trivial:

$$t_{ML}(f|e) = \frac{\text{count}(e, f)}{\text{count}(e)} \quad q_{ML}(j|i, l, m) = \frac{\text{count}(j, i, l, m)}{\text{count}(i, l, m)}$$

Input: A training corpus $(f^{(k)}, e^{(k)}, a^{(k)})$ for $k = 1 \dots n$, where $f^{(k)} = f_1^{(k)} \dots f_{m_k}^{(k)}$, $e^{(k)} = e_1^{(k)} \dots e_{l_k}^{(k)}$, $a^{(k)} = a_1^{(k)} \dots a_{m_k}^{(k)}$.

Algorithm:

- ▶ Set all counts $c(\dots) = 0$
- ▶ For $k = 1 \dots n$
 - ▶ For $i = 1 \dots m_k$, For $j = 0 \dots l_k$,

$$c(e_j^{(k)}, f_i^{(k)}) \leftarrow c(e_j^{(k)}, f_i^{(k)}) + \delta(k, i, j)$$

$$c(e_j^{(k)}) \leftarrow c(e_j^{(k)}) + \delta(k, i, j)$$

$$c(j|i, l, m) \leftarrow c(j|i, l, m) + \delta(k, i, j)$$

$$c(i, l, m) \leftarrow c(i, l, m) + \delta(k, i, j)$$

where $\delta(k, i, j) = 1$ if $a_i^{(k)} = j$, 0 otherwise.

Output: $t_{ML}(f|e) = \frac{c(e, f)}{c(e)}$, $q_{ML}(j|i, l, m) = \frac{c(j|i, l, m)}{c(i, l, m)}$

Parameter Estimation with the EM Algorithm

- Input: $(e^{(k)}, f^{(k)}), k = 1 \dots n$
Each $e^{(k)}$ is an English sentence, each $f^{(k)}$ is a French sentence
- The algorithm is related to algorithm with observed alignments, but with two key differences:
 - Iterative: start with initial (e.g., random) choice of q and t parameters, at each iteration: compute some “counts” base on data and parameters, and re-estimate parameters
 - The definition of of the delta function is different:

$$\delta(k, i, j) = \frac{q(j|i, l_k, m_k) t(f_i^{(k)} | e_j^{(k)})}{\sum_{j=0}^{l_k} q(j|i, l_k, m_k) t(f_i^{(k)} | e_j^{(k)})}$$

Input: A training corpus $(f^{(k)}, e^{(k)})$ for $k = 1 \dots n$, where $f^{(k)} = f_1^{(k)} \dots f_{m_k}^{(k)}$, $e^{(k)} = e_1^{(k)} \dots e_{l_k}^{(k)}$.

Initialization: Initialize $t(f|e)$ and $q(j|i, l, m)$ parameters (e.g., to random values).

For $s = 1 \dots S$

- ▶ Set all counts $c(\dots) = 0$
- ▶ For $k = 1 \dots n$
 - ▶ For $i = 1 \dots m_k$, For $j = 0 \dots l_k$

$$c(e_j^{(k)}, f_i^{(k)}) \leftarrow c(e_j^{(k)}, f_i^{(k)}) + \delta(k, i, j)$$

$$c(e_j^{(k)}) \leftarrow c(e_j^{(k)}) + \delta(k, i, j)$$

$$c(j|i, l, m) \leftarrow c(j|i, l, m) + \delta(k, i, j)$$

$$c(i, l, m) \leftarrow c(i, l, m) + \delta(k, i, j)$$

where

$$\delta(k, i, j) = \frac{q(j|i, l_k, m_k) t(f_i^{(k)} | e_j^{(k)})}{\sum_{j=0}^{l_k} q(j|i, l_k, m_k) t(f_i^{(k)} | e_j^{(k)})}$$

- ▶ Recalculate the parameters:

$$t(f|e) = \frac{c(e, f)}{c(e)} \quad q(j|i, l, m) = \frac{c(j|i, l, m)}{c(i, l, m)}$$

$$\delta(k, i, j) = \frac{q(j|i, l_k, m_k) t(f_i^{(k)} | e_j^{(k)})}{\sum_{j=0}^{l_k} q(j|i, l_k, m_k) t(f_i^{(k)} | e_j^{(k)})}$$

$e^{(100)}$ = And the program has been implemented

$f^{(100)}$ = Le programme a ete mis en application

For $s = 1 \dots S$

- ▶ Set all counts $c(\dots) = 0$
- ▶ For $k = 1 \dots n$
 - ▶ For $i = 1 \dots m_k$, For $j = 0 \dots l_k$

$$c(e_j^{(k)}, f_i^{(k)}) \leftarrow c(e_j^{(k)}, f_i^{(k)}) + \delta(k, i, j)$$

$$c(e_j^{(k)}) \leftarrow c(e_j^{(k)}) + \delta(k, i, j)$$

$$c(j|i, l, m) \leftarrow c(j|i, l, m) + \delta(k, i, j)$$

$$c(i, l, m) \leftarrow c(i, l, m) + \delta(k, i, j)$$

where

$$\delta(k, i, j) = \frac{q(j|i, l_k, m_k) t(f_i^{(k)} | e_j^{(k)})}{\sum_{j=0}^{l_k} q(j|i, l_k, m_k) t(f_i^{(k)} | e_j^{(k)})}$$

- ▶ Recalculate the parameters:

$$t(f|e) = \frac{c(e, f)}{c(e)} \quad q(j|i, l, m) = \frac{c(j|i, l, m)}{c(i, l, m)}$$

Justification for the Algorithm

- Input: $(e^{(k)}, f^{(k)}), k = 1 \dots n$

Each $e^{(k)}$ is an English sentence, each $f^{(k)}$ is a French sentence

- The log-likelihood function:

$$L(t, q) = \sum_{k=1}^n \log p(f^{(k)} | e^{(k)}) = \sum_{k=1}^n \log \sum_a p(f^{(k)}, a | e^{(k)})$$

- The maximum-likelihood estimates are:

$$\arg \max_{t, q} L(t, q)$$

- The EM algorithm will converge to a local maximum of the log-likelihood function

Summary

- Key ideas in the IBM translation models:
 - Alignment variables
 - Translation parameters, e.g., $t(\text{chien}|\text{dog})$
 - Distortion parameters, e.g., $q(2|1,6,7)$
- The EM algorithm: an iterative algorithm for training the q and t parameters
- Once parameters are trained, can recover the most likely alignment on our training examples

$e^{(100)}$ = And the program has been implemented

$f^{(100)}$ = Le programme a ete mis en application