

# Statistical Machine Translation: Word Based Translation Models

Michael Wohlmayr

# Statistical Machine Translation

- There is not THE ONE english translation **e** of a foreign sentence **f**.
- Some translations **e** are more likely than others:

*Das ist ein schnelles Auto.*

probable:

*It is a rapid vehicle.*

probable:

*This is a fast car.*

less probable:

*It's a sunny day.*

# Statistical Machine Translation

- Assign each sentence pair (**e**,**f**) a probability  $p(\mathbf{e}|\mathbf{f})$
- Assume we have a model to calculate  $p(\mathbf{e}|\mathbf{f})$
- To translate a foreign sentence into english, find:

$$\hat{\mathbf{e}} = \arg \max_{\mathbf{e}} p(\mathbf{e} | \mathbf{f})$$

# Well formed strings

- Problem:  $p(\mathbf{e}|\mathbf{f})$  must concentrate its probability on *well-formed* english sentences

| <i>well-formed</i>     | <i>ill-formed</i>      |
|------------------------|------------------------|
| I live in a house.     | I house a in live.     |
| Have you seen my keys? | Seen keys have my you? |

# Bayes Rule

- Relax constraint by using alternative statement:

$$\hat{\mathbf{e}} = \arg \max_{\mathbf{e}} p(\mathbf{e}) p(\mathbf{f} | \mathbf{e})$$

- Bayesian Reasoning: „What is the chance that the producer of  $\mathbf{f}$  had  $\mathbf{e}$  in mind and then translated it to  $\mathbf{f}$ “

*Source Channel Model*

# Language and Translation Model

$$\hat{\mathbf{e}} = \arg \max_{\mathbf{e}} p(\mathbf{e}) p(\mathbf{f} | \mathbf{e})$$

- $p(\mathbf{e})$ : language model. Gives a low probability to ill-formed strings.
- $p(\mathbf{f}|\mathbf{e})$ : translation model. Must not concentrate on well-formed strings anymore. Gives the probability that a english *bag of words* will translate into a french *bag of words*.
- *Probabilistic model split into 2 simpler models.*

# Decoding

$$\hat{\mathbf{e}} = \arg \max_{\mathbf{e}} p(\mathbf{e}) p(\mathbf{f} | \mathbf{e})$$

- Process of finding  $\mathbf{e}$  that maximizes the above product.

# Translation Model

- „Explains“ how english sentence becomes a french sentence.
- Based on a set of parameters.
- Given the parameters and a sentence pair ( $\mathbf{e}, \mathbf{f}$ ), computation of  $p(\mathbf{f}|\mathbf{e})$  is possible.
- IBM Model 1-5, developed in *Candide* Project
- Lets start with model 3...



# Model 3

- Very simplistic way of explaining how sentence **e** is translated to foreign sentence **f**
- Based on probabilities, simply reproduce each word a number of times, translate each word then reposition them
- Don't worry: This model is not used for translation, just to judge  $p(\mathbf{f}|\mathbf{e})$

# Fertility


- Each word in the english sentence may produce a certain amount of foreign words.
- e.g.

*This is not true.*

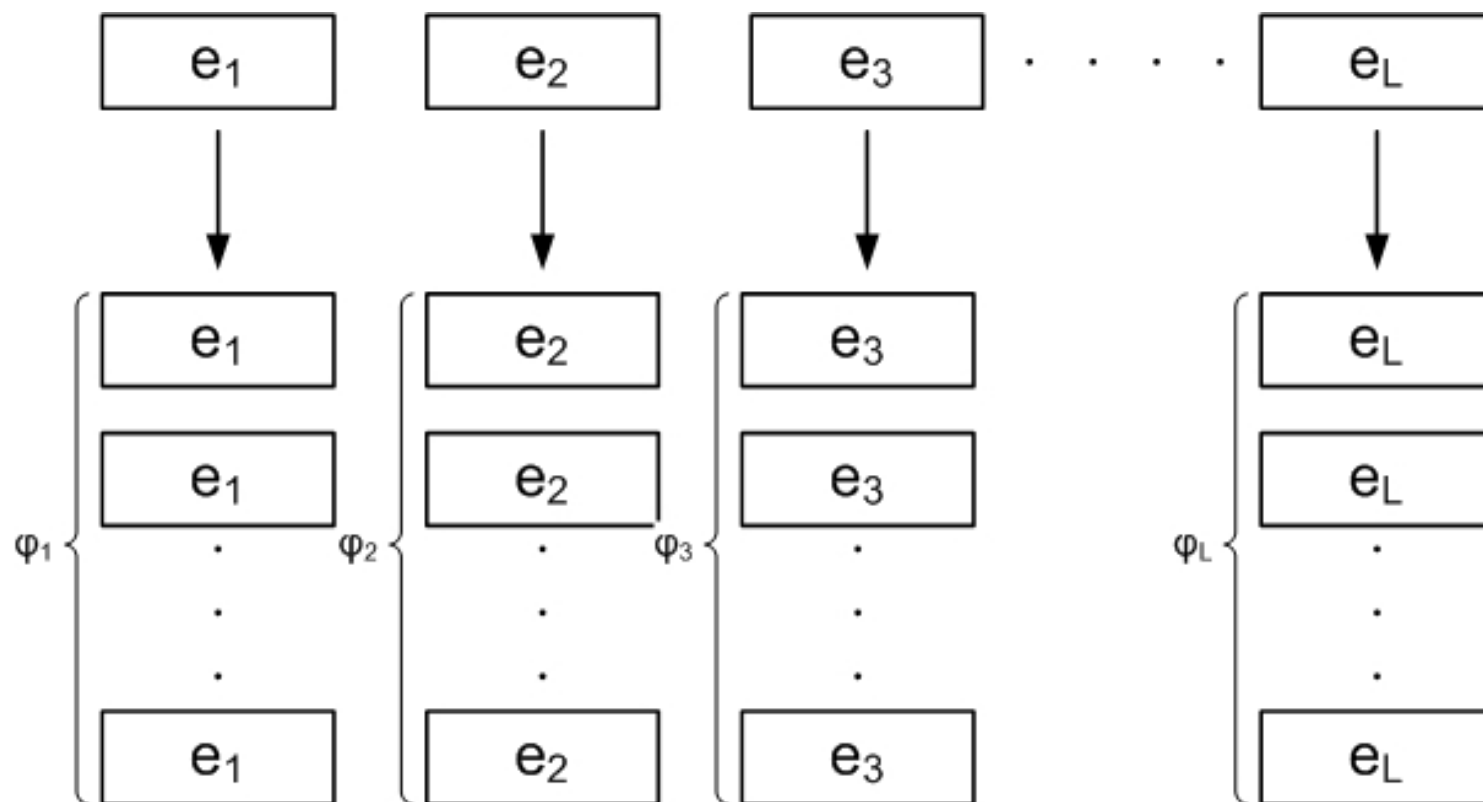
*Ce n'est pas vrais.*

# Fertility

- Fertility probability  $p_n(\phi|e)$

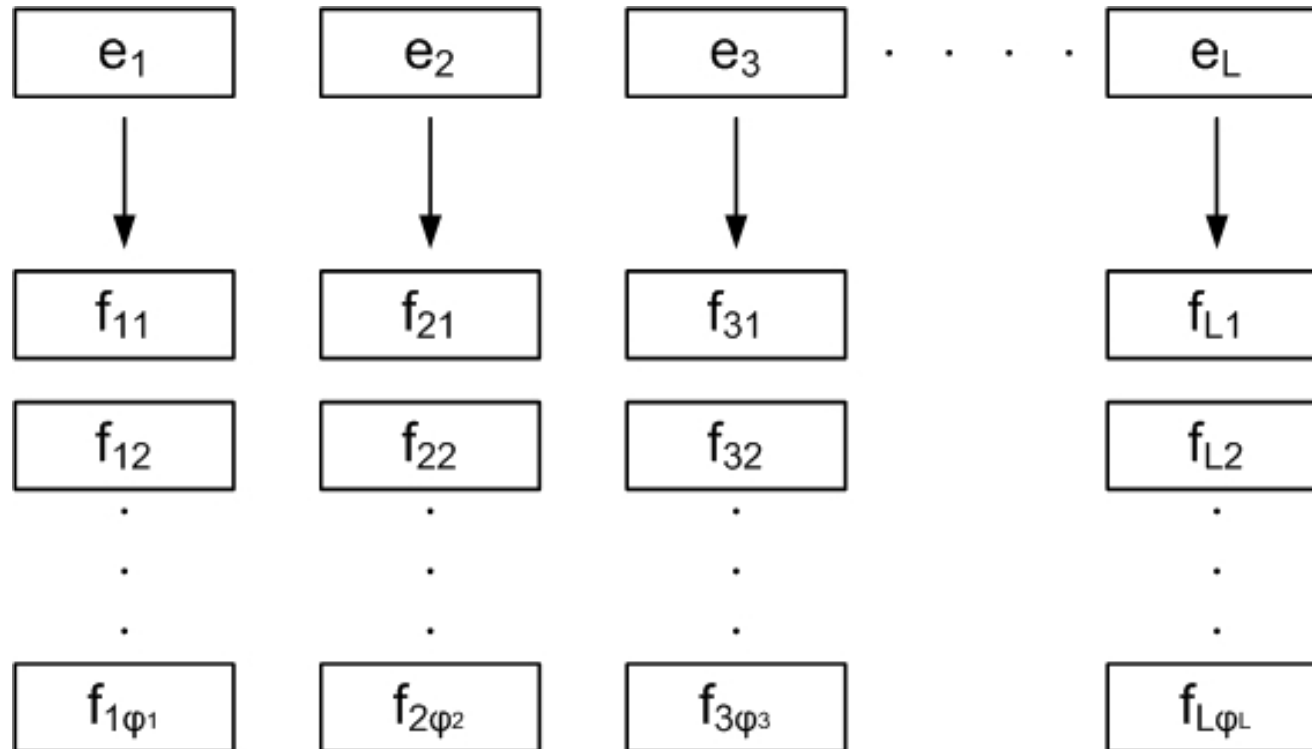


# Fertility



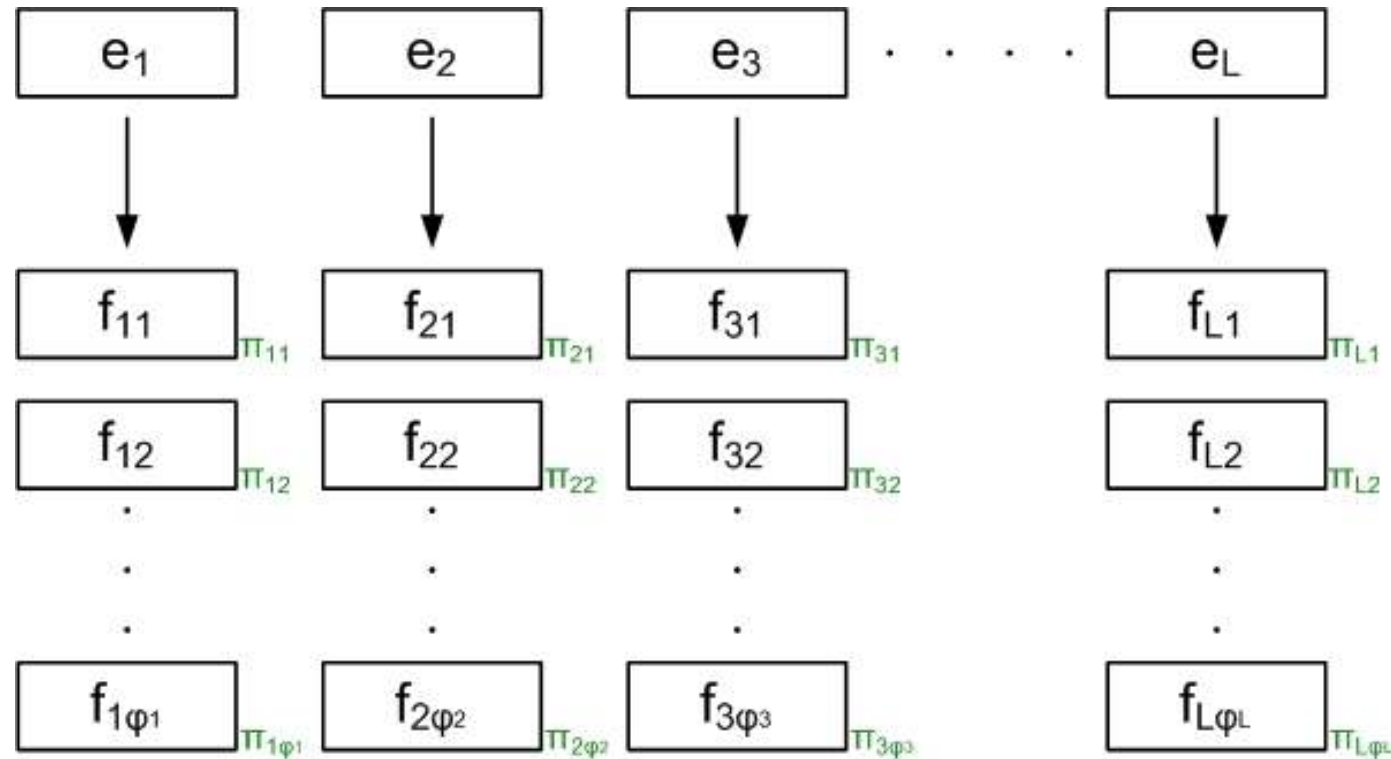
# Word Translation

- Word translation probability  $p_t(f|e)$

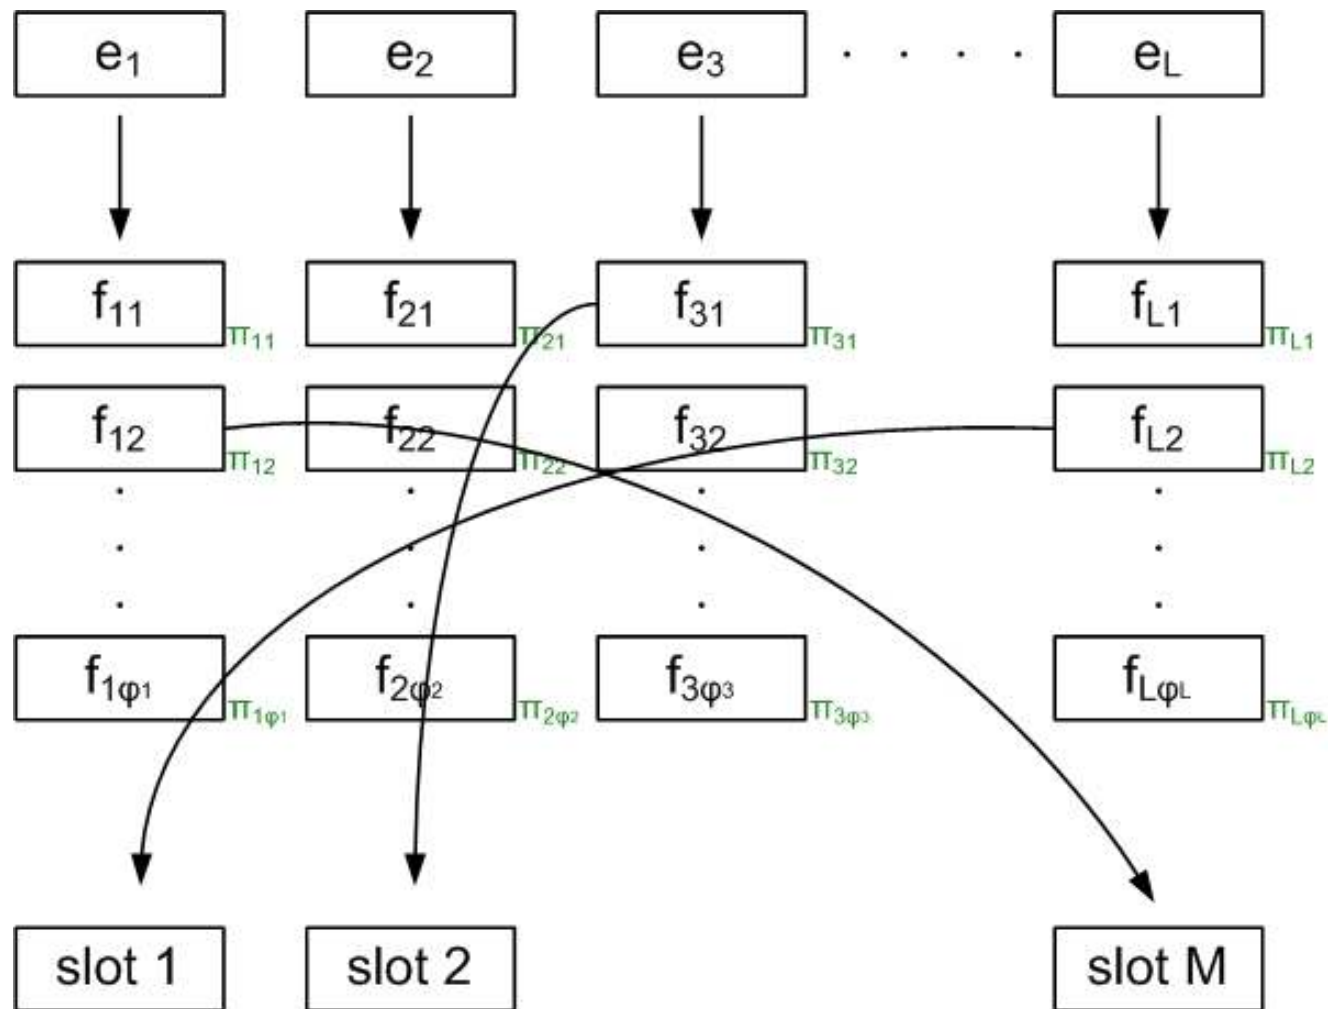


# Positioning

- distortion probability  $p_p(\pi|i,L,M)$



# Positioning



# Spurious Words

- Words that appear in the translation „although no word in the english sentence can be held responsible for them“

*I go home.*

*Ich gehe **nach** Hause.*



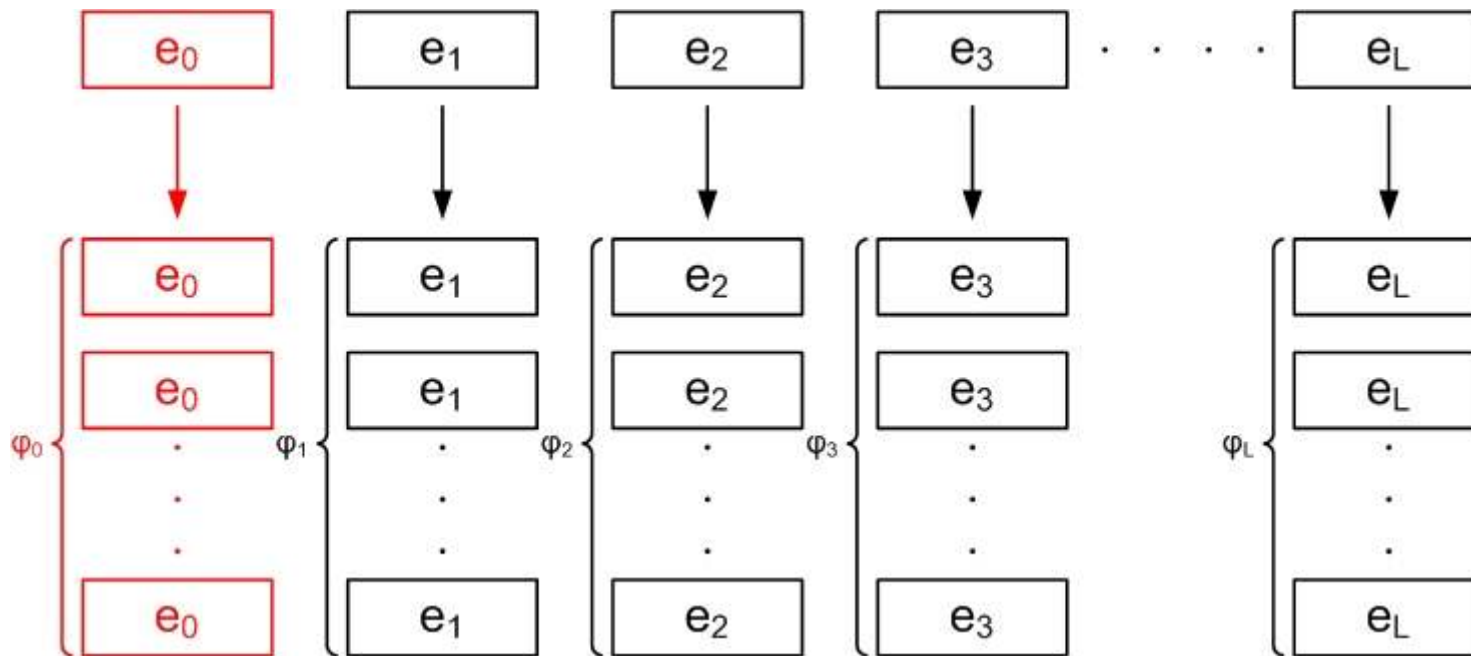
# Spurious Words

- Spurious words are added after the „normal“ words have been generated.
- After each foreign word, add a spurious word with probability  $p_1$

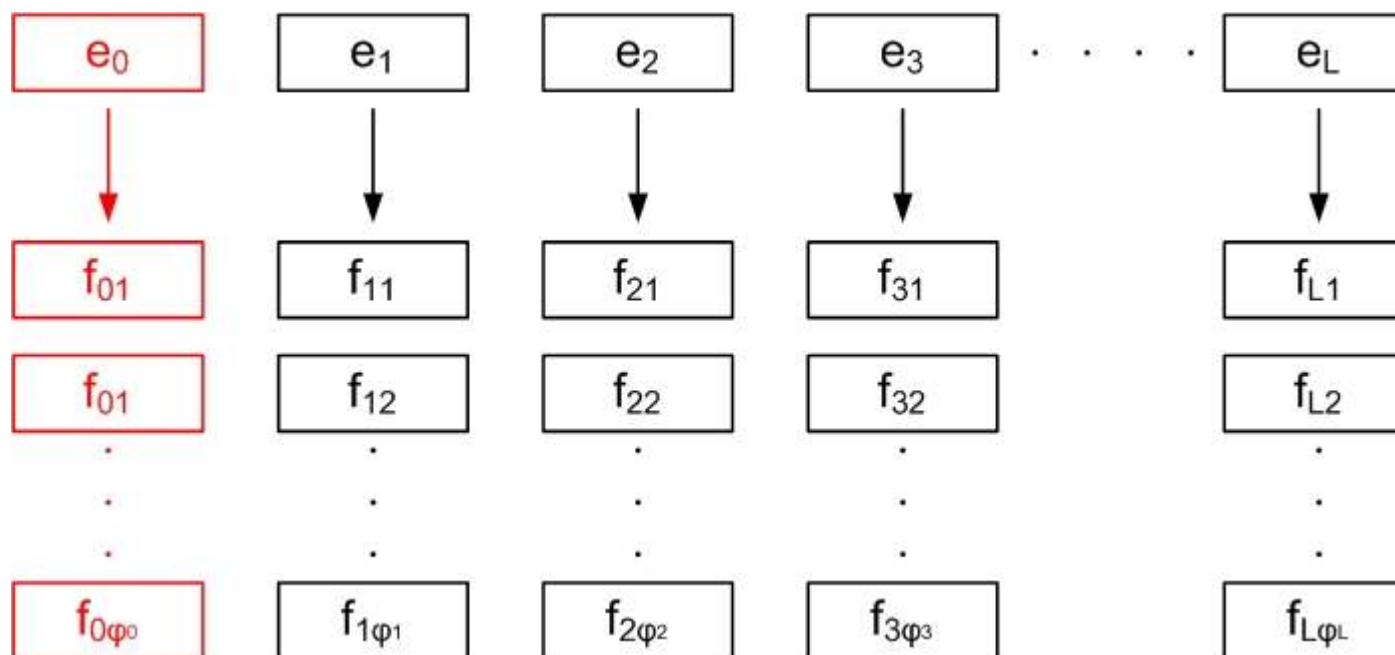
$$M_n = \varphi_1 + \varphi_2 + \dots + \varphi_L$$

$$p(\varphi_0) = \binom{M_n}{\varphi_0} (1 - p_1)^{M_n - \varphi_0} p_1^{\varphi_0}$$

# Spurious Words




# Word Translation




# Positioning of spurious words

- Final foreign sentence has  $M = M_n + \varphi_0$  words.
- After positioning the normal words, assign each spurious word a random position in one of the remaining free positions.

# Positioning of spurious words



# Positioning of spurious words



# Model 3 Parameters

- fertility  $p_n(\phi|e)$
- translation  $p_t(f|e)$
- positioning  $p_p(\pi|i,L,M)$
- spurious fertility  $p_1$

**HOW TO ESTIMATE?**

# Training sentences

- Based on large corpora, consisting of many translation pairs (e,f)
- e.g.: European Parliament Proceedings Parallel Corpus 1996-2003


<http://people.csail.mit.edu/koehn/publications/euoparl/>

Mr President , I respond to an invitation yesterday afternoon by the President of the House to speak on behalf of my group on a matter referred to in the Minutes .

Herr Präsident , ich entspreche hiermit einer vom Präsidenten des Hauses gestern nachmittag geäußerten Aufforderung , im Namen meiner Fraktion zu der im Protokoll genannten Angelegenheit Stellung zu nehmen .




# Word-for-Word Alignments




- Each word  $e$  might be connected to one or more words in  $f$
- Each word  $f$  is connected to one and only one word in  $e$ .

# Word-for-Word Alignments



- Given a large corpus together with alignment information, parameters can be estimated
- Count events in each sentence

# Fertility Estimation




$a = (1, 2, 0, 3)$

- $c_n(1|I) += 1$
- $c_n(1|go) += 1$
- $c_n(1|home) += 1$
- $c_n(1|NULL) += 1$

$$p_1 = \frac{\text{\#spurious words in translations}}{\text{\#normal words in translations}}$$


# Translation Estimation



$a = (1, 2, 0, 3)$

- $c_t(\text{Ich}|\text{I}) += 1$
- $c_t(\text{gehe}|\text{go}) += 1$
- $c_t(\text{hause}|\text{home}) += 1$
- $c_t(\text{nach}|\text{NULL}) += 1$

# Positioning Estimation



$a = (1, 2, 0, 3)$

- $c_p(1|1,3,4) += 1$
- $c_p(2|2,3,4) += 1$
- $c_p(4|3,3,4) += 1$

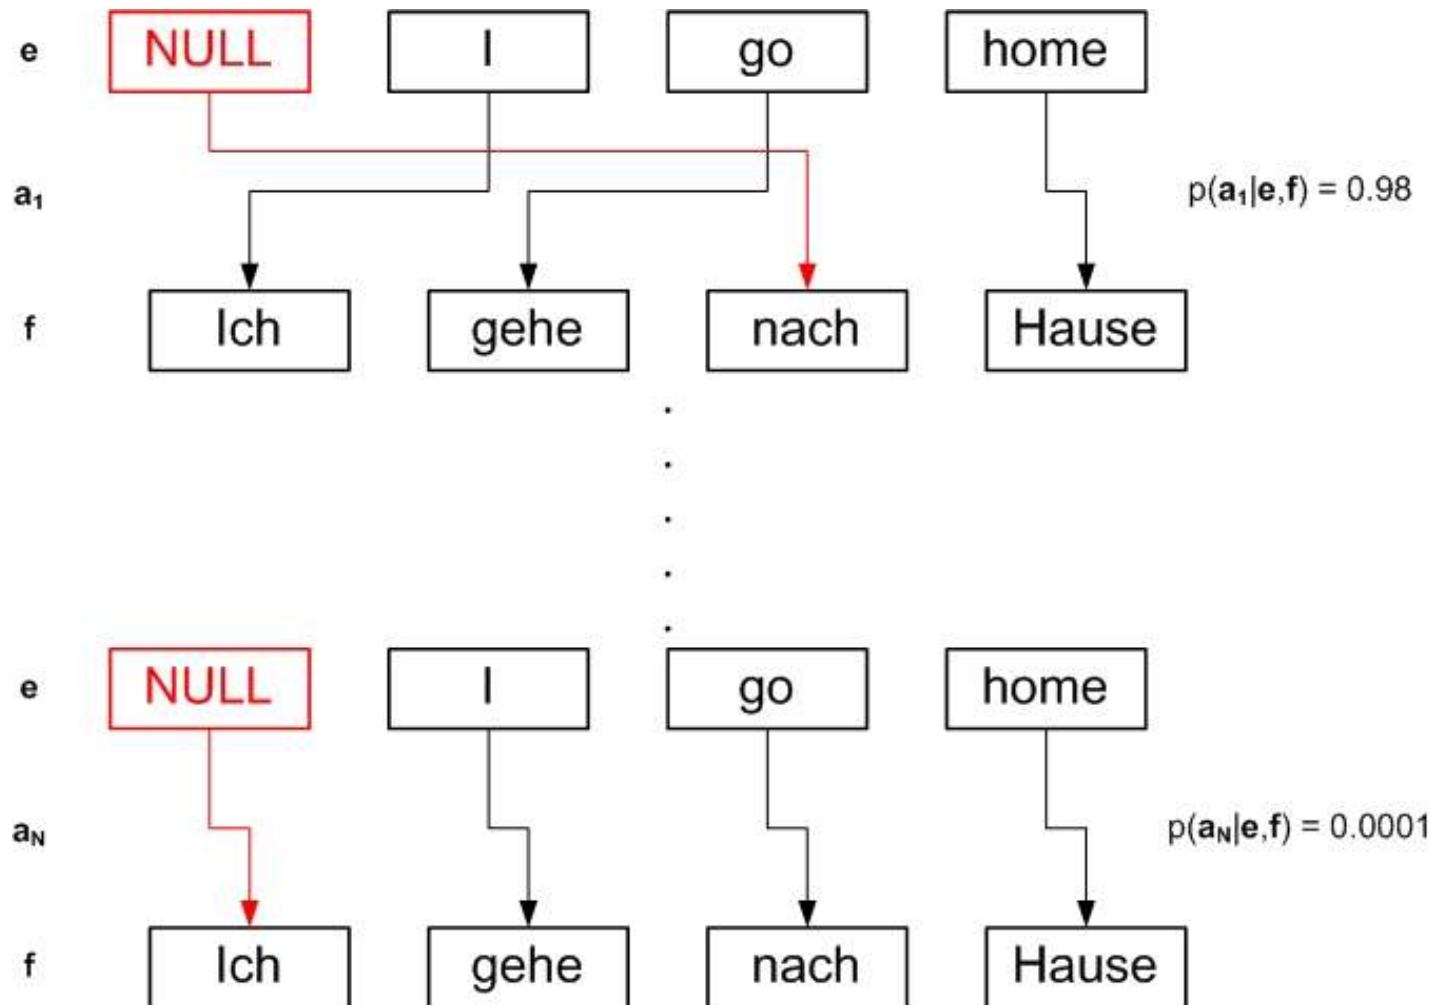
# Parameter Estimation

- After collecting these counts over the whole corpus, normalize them to form a proper probability distribution
- **PROBLEM:** We don't have alignment information

# Alignment Probability

- For each  $(\mathbf{e}, \mathbf{f})$ , there are  $(L+1)^M$  possible alignments
- Some alignments are more probable than others
- Alignment probability  $p(\mathbf{a}|\mathbf{e}, \mathbf{f})$

# Alignment Probability





# Parameter Estimation

$$c_n(\varphi | e_i) += \sum_{\mathbf{a}: \{e_i \leftrightarrow (f_1, f_2, \dots, f_\varphi)\}} p(\mathbf{a} | \mathbf{e}, \mathbf{f})$$

$$c_t(f_j | e_i) += \sum_{\mathbf{a}: \{e_i \leftrightarrow f_j\}} p(\mathbf{a} | \mathbf{e}, \mathbf{f})$$

$$c_p(j | i, L, M) += \sum_{\mathbf{a}: \{e_i \leftrightarrow f_j\}} p(\mathbf{a} | \mathbf{e}, \mathbf{f})$$

# Parameter Estimation

- With alignment probabilities, parameter estimation is possible too.
- Again the same problem: How do we get alignment probabilities?

# Alignment Probability

- Reformulate:

$$p(\mathbf{a} | \mathbf{e}, \mathbf{f}) = \frac{p(\mathbf{a}, \mathbf{f} | \mathbf{e})}{p(\mathbf{f} | \mathbf{e})} = \frac{p(\mathbf{a}, \mathbf{f} | \mathbf{e})}{\sum_{\mathbf{a}} p(\mathbf{a}, \mathbf{f} | \mathbf{e})}$$

- With  $p(\mathbf{a}, \mathbf{f} | \mathbf{e})$ , we can compute:
  - alignment probabilities  $p(\mathbf{a} | \mathbf{e}, \mathbf{f})$
  - $p(\mathbf{f} | \mathbf{e})$

$$p(\mathbf{a}, \mathbf{f} | \mathbf{e})$$

- ...the probability that a certain translation  $\mathbf{f}$  with alignment  $\mathbf{a}$  will be generated, given  $\mathbf{e}$ .
- Assume we have all model parameters. Then we can compute  $p(\mathbf{a}, \mathbf{f} | \mathbf{e})$

$$p(\mathbf{a}, \mathbf{f} | \mathbf{e})$$

- Approximation:

$$p(\mathbf{a}, \mathbf{f} | \mathbf{e}) \approx \prod_{i=1}^L p_n(\varphi_i | e_i) \prod_{j=1}^M p_t(f_j | e_{a(j)}) \prod_{\{j|a(j) \neq 0\}} p_p(j | a(j), L, M)$$

- Still missing:
  - probability for  $\varphi_0$
  - $\varphi_0!$  ways to arrange spurious words in free slots
  - $\varphi_i!$  ways to end up in alignment  $e_i \leftrightarrow (f_1, f_2, \dots, f_{\varphi_i})$

$$p(\mathbf{a}, \mathbf{f} | \mathbf{e})$$

- Finally:

$$p(\mathbf{a}, \mathbf{f} | \mathbf{e}) = \prod_{i=1}^L p_n(\varphi_i | e_i) \prod_{j=1}^M p_t(f_j | e_{a(j)}) \prod_{j:a(j) \neq 0} p_p(j | a(j), L, M) \\ \binom{M - \varphi_0}{\varphi_0} (1 - p_1)^{M - 2\varphi_0} p_1^{\varphi_0} \prod_{i=0}^L \varphi_i! \frac{1}{\varphi_0!}$$

- ...depends on model parameters only.

# And now???

- To estimate the model parameters  $\theta$ , we need  $p(\mathbf{a}, \mathbf{f} | \mathbf{e})$
- To estimate  $p(\mathbf{a}, \mathbf{f} | \mathbf{e})$ , we need the model parameters  $\theta$

# Expectation Maximization

- The EM-Algorithm is an iterative procedure to compute the Maximum Likelihood (ML) estimate in the presence of hidden or missing data
- We search the ML for:

$$p(\mathbf{f} \mid \theta, \mathbf{e}) = \sum_{\mathbf{a}} p(\mathbf{a}, \mathbf{f} \mid \theta, \mathbf{e})$$

- Hidden data  $\mathbf{a}$



# Expectation Maximization

$$\theta = \arg \max_{\theta} \sum_{\mathbf{a}} p(\mathbf{a} \mid \mathbf{f}, \theta_n, \mathbf{e}) \ln p(\mathbf{a}, \mathbf{f} \mid \theta, \mathbf{e})$$

- Expectation: Estimate  $p(\mathbf{a} \mid \mathbf{f}, \mathbf{e})$  using model parameters  $\theta_n$  from the previous iteration
- Maximization: Estimate optimal model parameters  $\theta_{n+1}$  using the estimated alignment probabilities

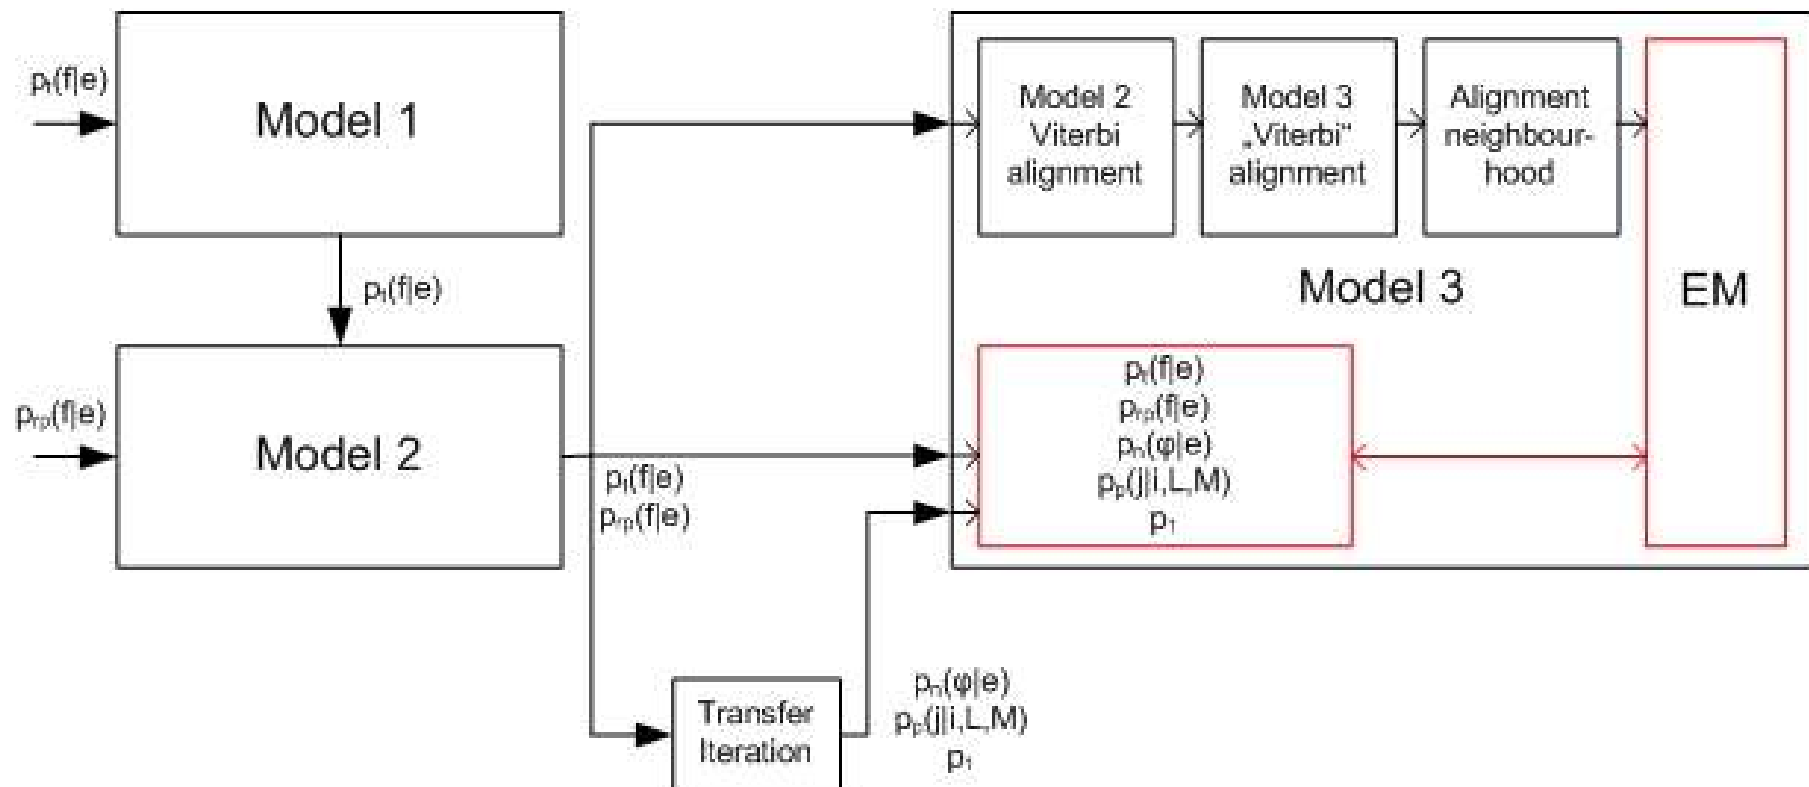
# Problems with Model 3 Training

- EM depends on parameter initialization and will find local optimum only
- EM needs to iterate over all possible alignments of a sentence pair. A sentence pair consisting each of 20 words has  $(21)^{20} = 2.7822\text{e}+026$  possible alignments

# Solution for Model 3 Training

- Initialize  $\theta_1$  with a *good guess*
- Don't iterate over all possible alignments, but only over a *subset of likely alignments*
- Use simpler models (model 1 and 2) to get these guesses

# Solution for Model 3 Training



# Model 1

- Forget about fertility and distortion:

$$p(\mathbf{a}, \mathbf{f} \mid \mathbf{e}) = \prod_{j=1}^M p_t(f_j \mid e_{a(j)})$$

- Computational advantage:

$$p(\mathbf{f} \mid \mathbf{e}) = \sum_{\mathbf{a}} \prod_{j=1}^M p_t(f_j \mid e_{a(j)}) = \prod_{j=1}^M \sum_{i=0}^L p_t(f_j \mid e_i)$$

# Model 1

- Using the same trick,  $p_t(f|e)$  can be trained with EM without iterating over all possible alignments

# Model 2

- Employ only translation and *reverse positioning* probability:

$$p(\mathbf{a}, \mathbf{f} \mid \mathbf{e}) = \prod_{j=1}^M p_t(f_j \mid e_{a(j)}) \prod_{j=1}^M p_{rp}(\mathbf{a}(j) \mid j, L, M)$$

- Again, we can use the same trick:

$$p(\mathbf{f} \mid \mathbf{e}) = \sum_{\mathbf{a}} \prod_{j=1}^M p_t(f_j \mid e_{a(j)}) p_{rp}(\mathbf{a}(j) \mid j, L, M) =$$

$$\prod_{j=1}^M \sum_{i=0}^L p_t(f_j \mid e_i) p_{rp}(i \mid j, L, M)$$

# Model 2

- Efficient estimation of  $p_t$  and  $p_{rp}$  possible
- Most probable (Viterbi) alignment of a sentence:

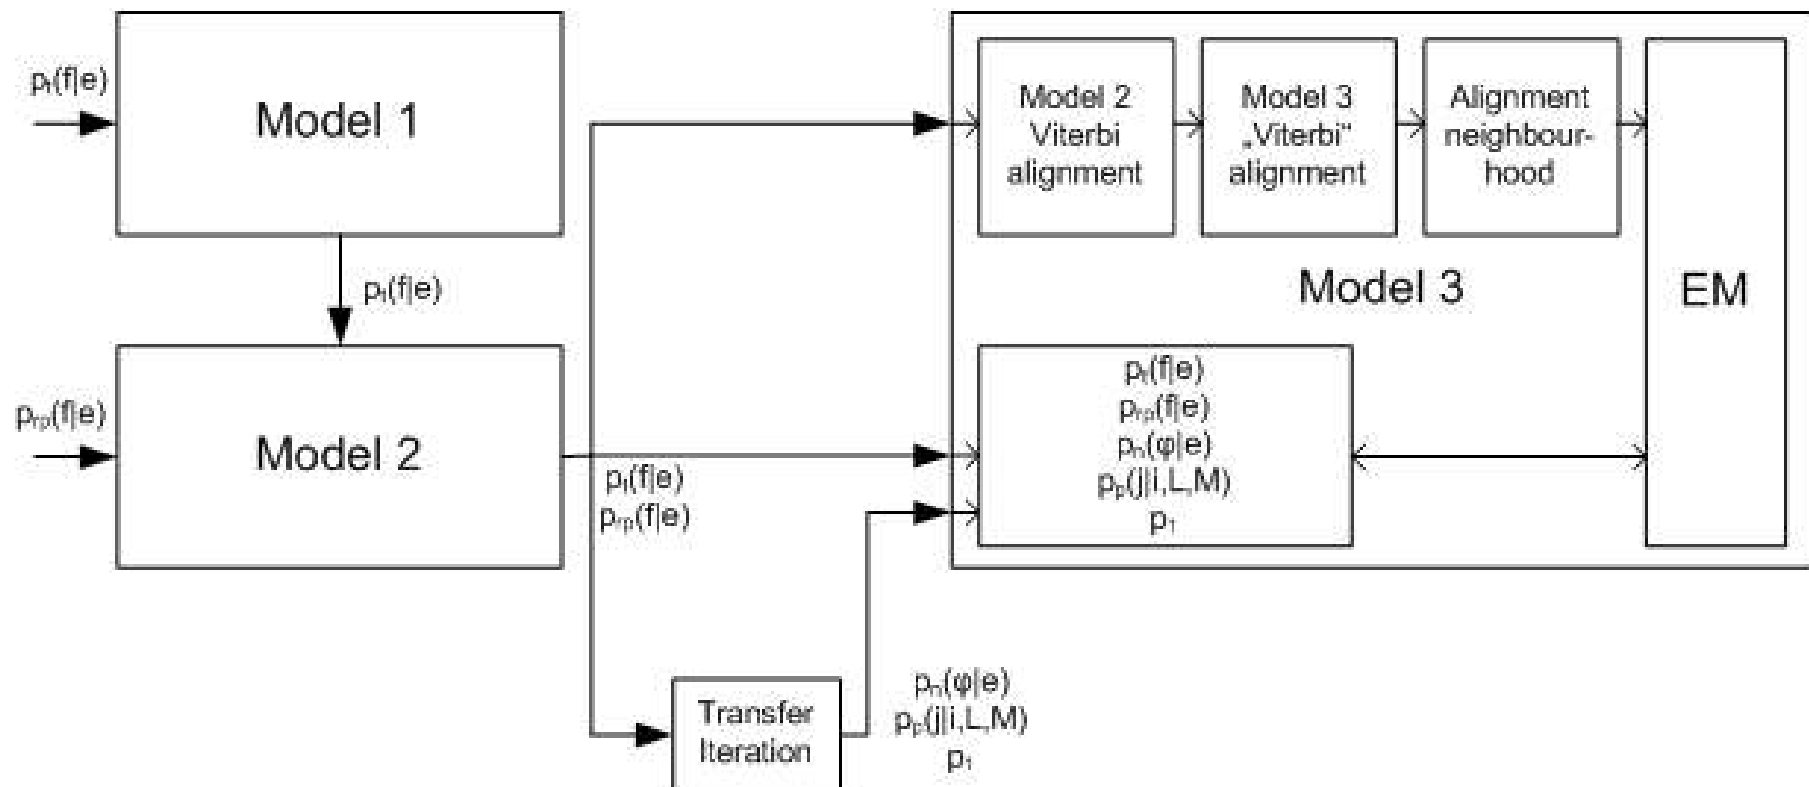
$$\mathbf{a}(j) = \arg \max_i p_t(f_j | e_i) p_{rp}(i | j, L, M) \quad \forall j$$



# Model 3 „Viterbi“ Alignment

- Model 2 Viterbi Alignment might be suboptimal in the sense of model 3
- Hillclimbing: iterative and quick method to find a better alignment. Replace  $\mathbf{a}$  with neighbour  $\mathbf{b}(\mathbf{a})$  which maximizes  $p(\mathbf{b}(\mathbf{a})|\mathbf{e},\mathbf{f})$ .
- Solution might still be suboptimal
- e.g.:
  - $\mathbf{a} = (1, 2, 3)$
  - $\mathbf{b}(\mathbf{a}) = (1, 2, 1)$  ...*differs by one move*
  - $\mathbf{b}(\mathbf{a}) = (1, 3, 2)$  ...*differs by one swap*

# IBM Model 1-3



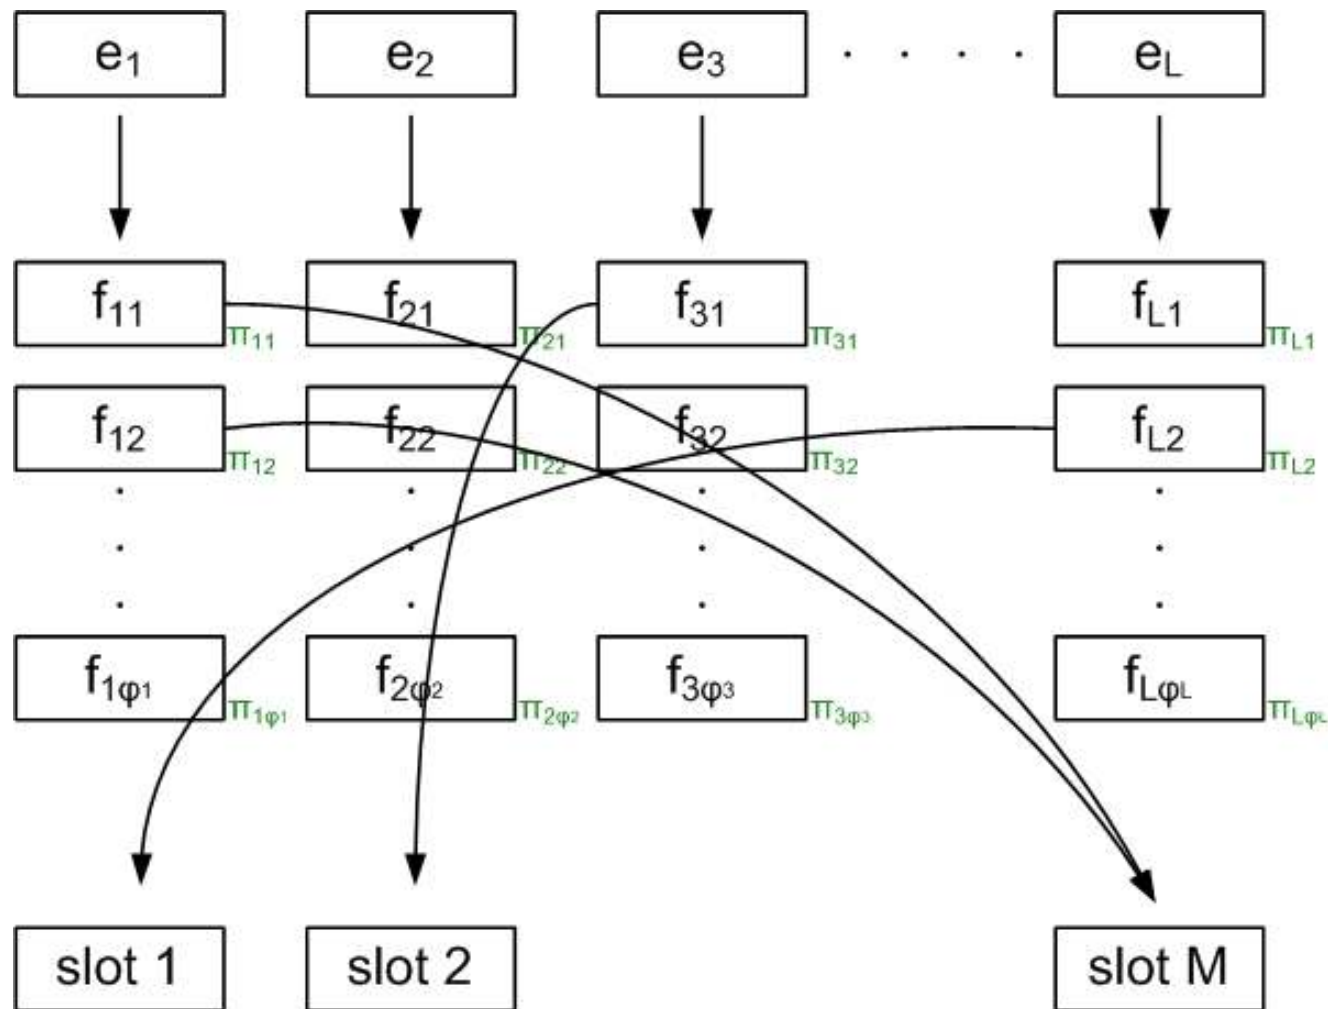
## 2 More Problems with Model 3

- $p_p(\pi|i,L,M)$  does not depend on words, just on positions. *Unrealistic*
- $p_p(\pi|i,L,M)$  has no memory. *Deficiency*

# Deficiency


- $p_p(\pi|i,L,M)$  does not consider values assigned to earlier words  $\rightarrow$  multiple words might get same position
- $p(\mathbf{a},\mathbf{f}|\mathbf{e})$  of Model 3 wastes some of its probability mass on impossible events, i.e. generalized strings

# Deficiency

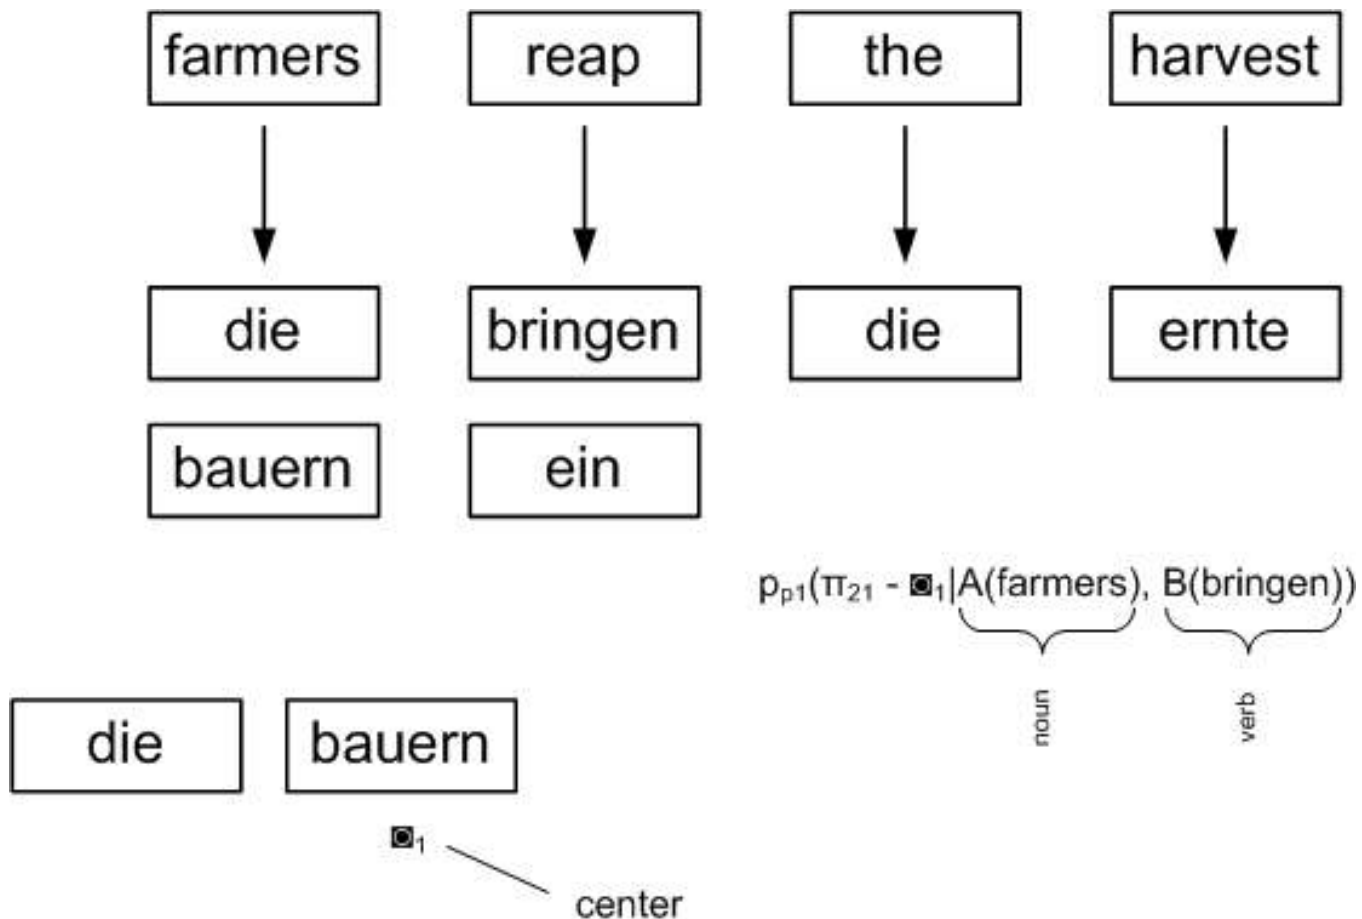


# Model 4


- Class based positioning of words
- e.g.:  $A_1$  verb,  $A_2$  noun,  $A_3$  adjective etc...



# Model4 Positioning



# Model4 Positioning



$$p_{\rho>1}(\pi_{22} - \pi_{21} | B(\text{ein}))$$

die    bauern    bringen



# Model 5


- Similar to Model 4, but prevents generation of invalid strings

# Conclusion

- Model 1-5 provide effective means for obtaining word-by-word alignments of translation

# Example

- Best out of  $5.6 \times 10^{31}$  alignments



# Conclusion

- Model 1-5 provide effective means for obtaining word-by-word alignments of translation
- Single parent constraint in alignments too simplistic
- Relation between several grammatical forms of a word is ignored.

# Example

- Different conjugations of one verb are treated separately

*should*

| f         | $t(f   e)$ | $\phi$ | $n(\phi   e)$ |
|-----------|------------|--------|---------------|
| devrait   | 0.330      | 1      | 0.649         |
| devraient | 0.123      | 0      | 0.336         |
| devrions  | 0.109      | 2      | 0.014         |
| faudrait  | 0.073      |        |               |
| faut      | 0.058      |        |               |
| doit      | 0.058      |        |               |
| aurait    | 0.041      |        |               |
| doivent   | 0.024      |        |               |
| devons    | 0.017      |        |               |
| devrais   | 0.013      |        |               |

# References

- Berger, Brown et al., *The Candide System for Machine Translation*, 1994
- Knight, *A Statistical MT Tutorial Workbook*, 1999
- Brown, Pietra et al., *The Mathematics of Statistical Machine Translation: Parameter Estimation*, 1993
- Borman, *The Expectation Maximization Algorithm - A Short Tutorial*, 2004